OPUS 5 · JULY 2026

Making the Most of Claude Opus 5

Dustin · 30 minutes, one live demo, three acts

FOR SOFTWARE ENGINEERS — PLUS PLATFORM ENGINEERS & DESIGNERS

FIVE DAYS AGO

It doubled its predecessor's hardest-benchmark score. The price didn't move.

Opus 5 more than doubles Opus 4.8's Frontier-Bench score — at the same **$5 / $25** per MTok, and at a lower cost *per task*.

COLD OPEN · TOLD, NOT SHOWN

A drawing of a machine part it had no way to look at. So it built its own eyes.

THE MID-2026 LINEUP IN 60 SECONDS

Four models, one dial each.

MYTHOS CLASS

Fable 5

$10 / $50 per MTok

Sits above Opus. Keeps SWE-bench Pro, factual breadth, creative polish.

THE CEILING

JUL 24 · NEW

Opus 5

$5 / $25 unchanged from 4.8

Near-Fable intelligence at half Fable's price. Now the default Opus in Claude Code.

THE WORKHORSE

JUN 30

Sonnet 5

$2 / $10 intro to Aug 31, then $3/$15

Most agentic Sonnet yet — performs close to Opus 4.8.

THE IMPLEMENTER

LEGACY · LOAD-BEARING

Haiku 4.5

Efficiency tier

Parallel execution, subagents, latency-sensitive paths.

THE FAN-OUT

Last month's flagship capability is now the mid-tier price.

OPUS 5 VS OPUS 4.8

What actually changed.

EFFORT · NOW A 5-LEVEL DIAL

`low` `med` `high` `xhigh` `max`

Set via `output_config.effort`. **high** is the API default — and the levels were recalibrated, so 4.8's settings don't carry over.

THINKING

On by default

Can only be disabled at effort ≤ high. Thinking-on at *low* usually beats thinking-off at similar cost.

CONTEXT

1M tokens

Default *and* max. Behavior stays consistent all the way through the window.

SPEED

Fast mode

~2.5× the speed at 2× the price. Interactive paths only — not your batch jobs.

BENCHMARKS

New SOTA

Frontier-Bench and GDPval-AA. Within **0.5%** of Fable 5's CursorBench peak at half the cost.

BEHAVIOR

Verifies itself

Checks its own work unprompted, iterates until it succeeds, pushes back on flawed designs. Review accuracy holds even at low effort.

FROM THE FIELD

Two stories, one behavior.

STORY A · THE PACKAGE-MANAGER BUG

It found the root cause. The other model found the symptom.

Opus 5 fixed an edge case the community's own patch had missed.

A competing model fixed the surface symptom — and reported it resolved.

STORY B · THE TRADING-FIRM FEED

No feed to test against, so it built one.

A market data feed for a new exchange, in one session.

With no live feed to validate against, it wrote its own test harness to check its parsing.

Both stories are the same story: it looks for the verifier.

THE HONEST COUNTERWEIGHT

A stronger builder than reviewer.

CodeRabbit's early review of Opus 5, running it across real review workloads.

WHAT GOT BETTER

Much better **design judgment** than 4.8.

Noticeably **"less anxious"** — fewer hedges, fewer defensive re-reads.

WHAT IT COSTS

~50%

MORE READ
PER REVIEW CALL

~65%

MORE WRITTEN
PER REVIEW CALL

Verbosity is now a line item. Which is why the fix is a delete list, not an add list.

THE MAP · COST VS OUTPUT

Effort moves you along a curve. Model choice jumps between curves.

Each dot is one (model, effort) pair. Hover any dot for the claim behind its position.

Haiku 4.5 · Sonnet 5 · Opus 4.8 (ref) · Opus 5 · Fable 5

show other labs

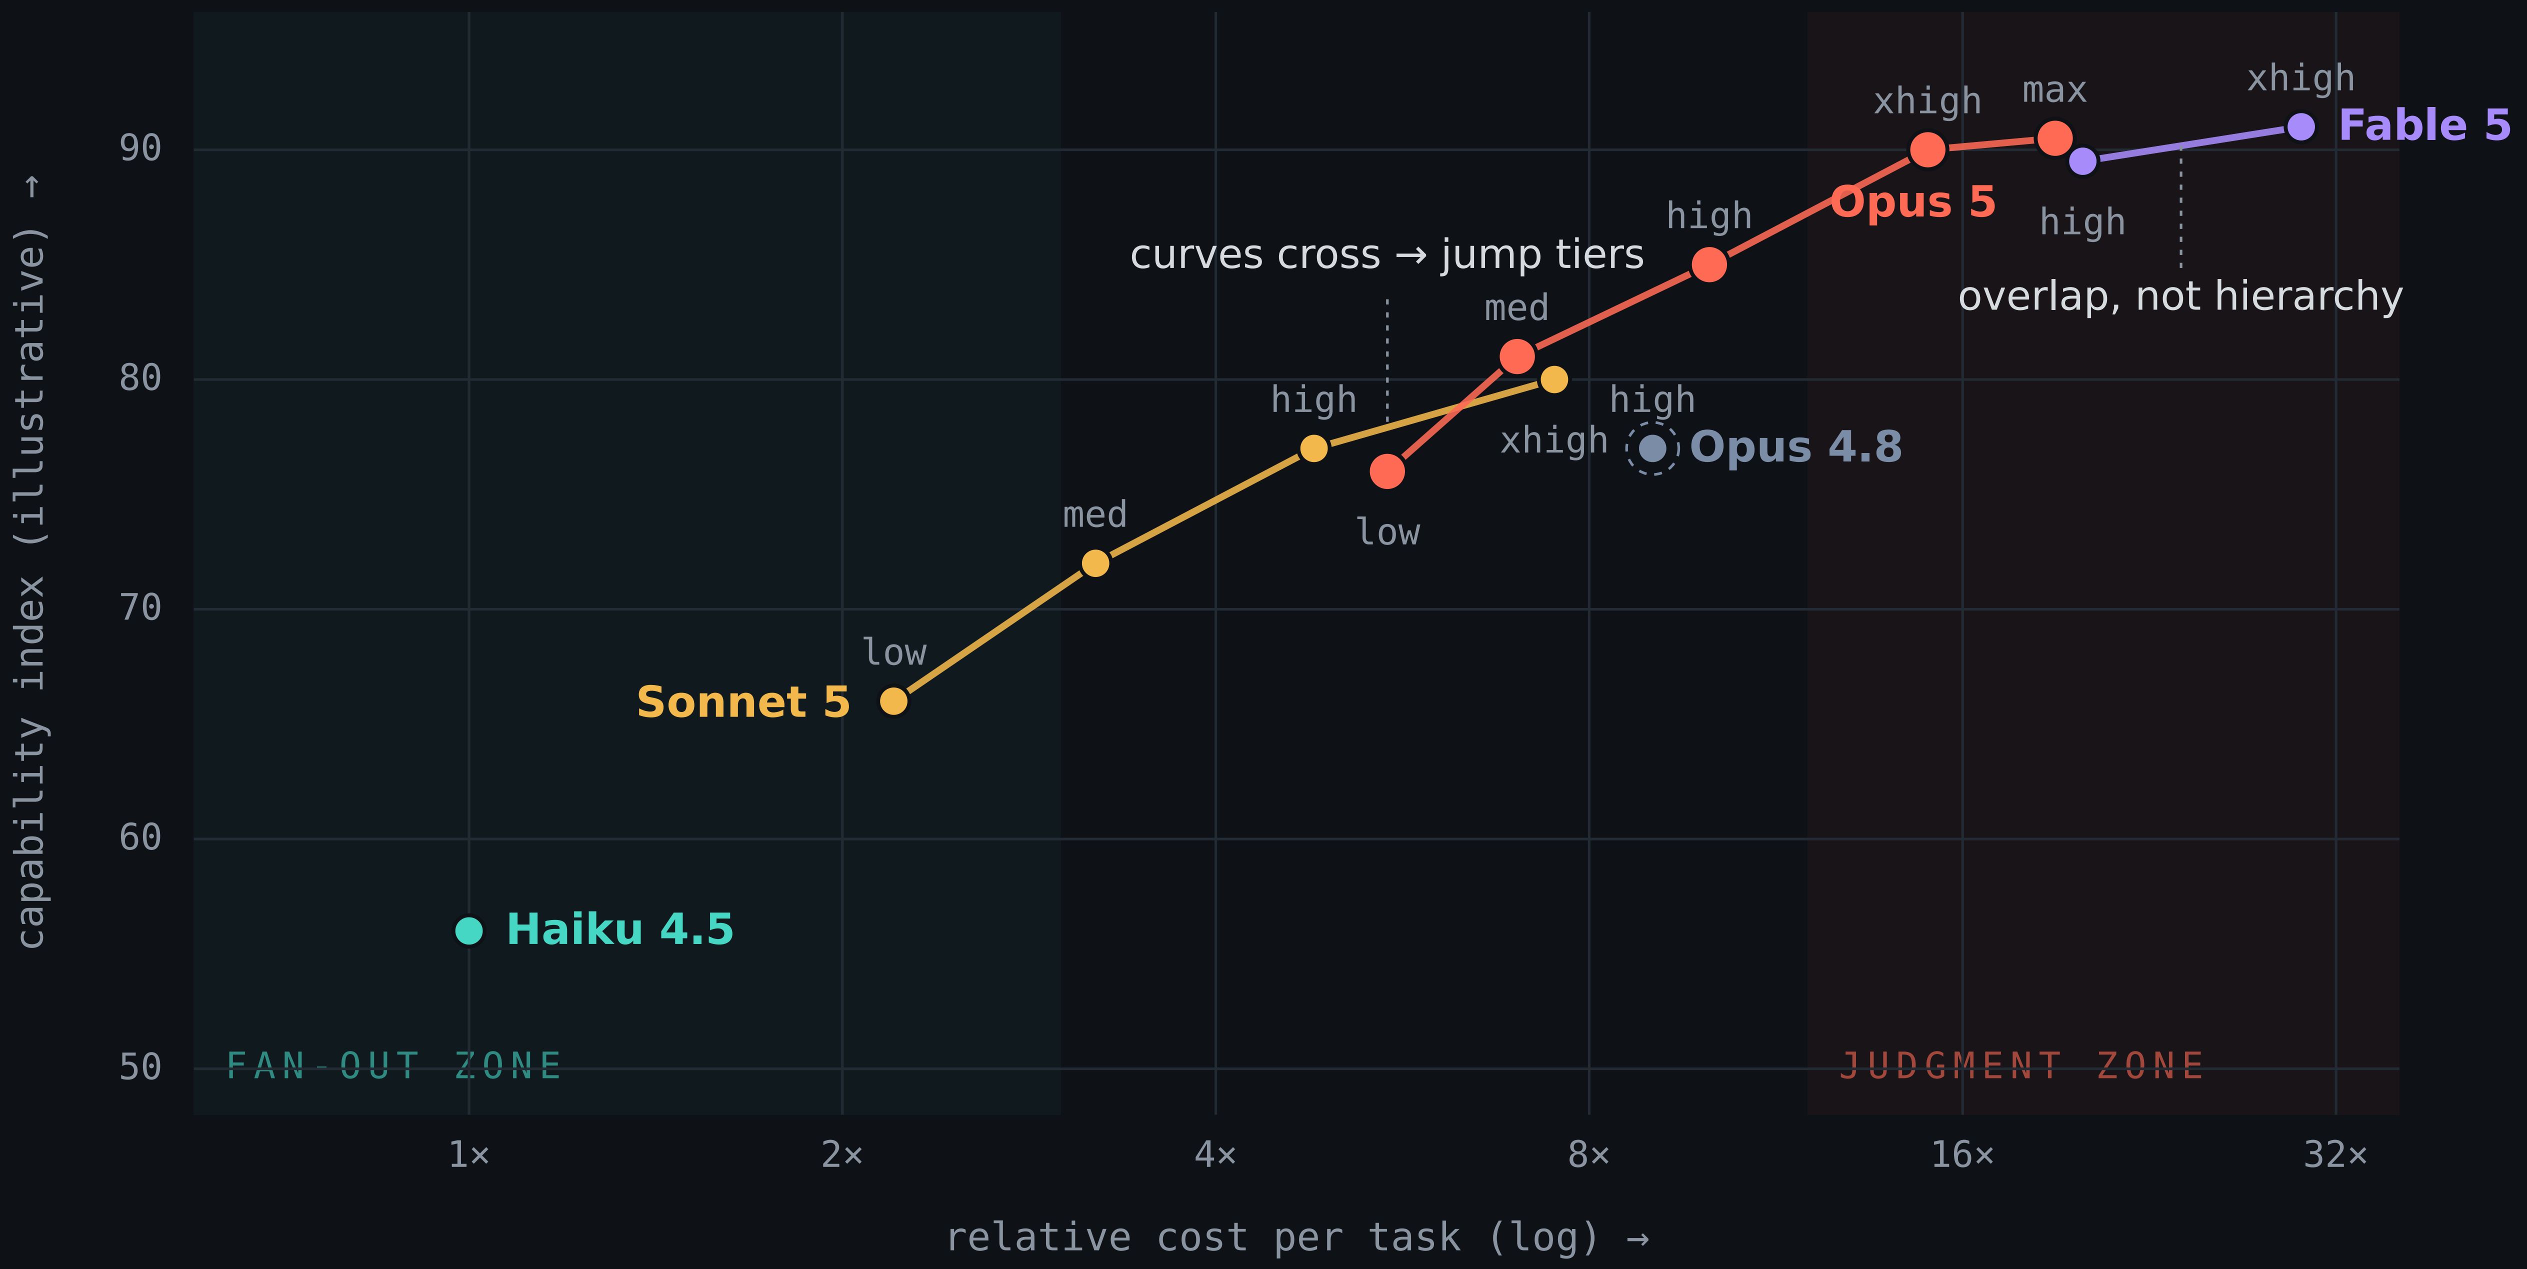

Positions are **illustrative** — synthesized from published pricing, Anthropic benchmark claims, and Artificial Analysis Intelligence-vs-Cost data. Per-effort scores aren't published.

tune effort before you switch models

THE TWO ZONES

The fan-out zone is priced per attempt. The judgment zone is priced per mistake.

FAN-OUT ZONE · LEFT, TEAL

Volume is the strategy.

Haiku, DeepSeek-class, Sonnet at low effort.

Mechanical, parallelizable work: read a hundred files, grep, summarize, extract.

Defining property: you run many at once and can afford to waste some output — an irrelevant file read costs a fraction of a cent.

= the model: haiku explorer, the mechanical nodes of a workflow

JUDGMENT ZONE · RIGHT, CORAL

Being wrong is the cost.

Opus 5 at high / xhigh, Fable.

Root cause, architecture, interpreting a hundred subagent reports, final review.

Defining property: you run it once — a misdiagnosed root cause costs days, so 15–20× per task is good economics.

= the orchestrator's seat

Tune effort before you switch models. Switch models when the curves cross.

The unshaded middle — Sonnet 5, Opus 5 low/med — is deliberate: negotiable territory where curve-crossing and effort sweeps decide.

DON'T TAKE MY WORD FOR IT

Every lab shipped a dial in 2026.

ARTIFICIAL ANALYSIS

This map, at industry scale.

Intelligence Index vs Cost per Task — Pareto-frontier framing: **everything off the frontier is paying more for less**. Now literally on SF billboards.

They plot effort levels as curves too: for any GPT-5.6 Terra effort level, a Luna or Sol level beats it. Same "curves cross" logic.

Broader context: 9 of the 13 models on their price-intelligence frontier are open-weights. The fan-out zone is contested territory.

THE CONVERGENCE

OpenAI	max reasoning + "ultra" parallel mode
Gemini	thinking levels + Deep Think
Grok	an effort parameter
Anthropic	low · med · high · xhigh · max

"Which model?" became **"which model, at what effort?"**

~5–10× / year **Pareto-frontier price-performance improvement. This map is a snapshot of a curve sliding down-and-right, fast.**

Caveat: cross-lab dials aren't apples-to-apples — comparisons normalize by token spend, not parameter names. Several headline numbers are vendor-reported.

MIGRATION CHECKLIST

Delete these from your prompts.

× ~~"Double-check your answer" · "verify with a subagent"~~
Over-verification. Wasted tokens, no quality gain — it already verifies.

× ~~Legacy harness verification steps~~
Same problem, wearing a config file.

× ~~"Only report high-severity issues"~~
It obeys literally and under-reports. Ask for everything, filter in a second pass.

× ~~Effort defaults carried over from 4.8~~
The levels were recalibrated. Re-run a sweep on your own evals.

✓ **ADD instead: conciseness instructions · scope constraints · delegation caps · narration cadence**
The four things that got *more* valuable, not less.

The model got more agentic, so your prompts should get less paranoid.

"Every six months, delete your CLAUDE.md, delete your skills, delete your hooks — see what the model does, and it might surprise you."

BORIS CHERNY · CREATOR OF CLAUDE CODE · YC 2026

MIGRATION · EVERY MAJOR MODEL RELEASE

The trial period: re-derive the harness.

1 **Archive, then delete**
Move `CLAUDE.md · skills/ · hooks/` into an archive folder, so nothing is lost — then start the next session with **none of it**.

2 **Use it bare for a week — don't guess**
"You don't want to guess what's the instruction that the model needs, because **you might not predict it correctly**." Run your real work. Watch.

3 **Add a line back only after the *same* stumble, repeatedly**
"Only when you see it repeatedly stumble on the same thing — that's when you add it back." One stumble is noise. **Three is a rule.**

4 **Charge rent per line**
"The model is going to read this instruction every single time you use it." Every line costs context on every turn. A config accumulating since Sonnet 3.5 isn't memory — it's **scar tissue from models that no longer exist**.

5 **Keep evals — expect them to saturate**
Code and prompts get deleted; evals outlive them by "maybe one, two, three model generations." Then the model saturates them, and you write new ones from where it now struggles.

The harness is something you re-derive per model — not an asset you maintain forever.

Boris Cherny — creator of Claude Code · YC Startup School 2026.

INTRODUCTION TO DYNAMIC WORKFLOWS

Delegation instincts, as infrastructure.

01 · you
Describe a complex task.

→

02 · claude
Writes a **JavaScript orchestration script**.

→

03 · runtime
Executes it in the background, fanning work out across subagents.

→

04 · back to you
Only the **coordinated result** returns.

Intermediate results stay in script variables — the coordination leaves your conversation, so your context stays clean and the fan-out runs in parallel.

FOR
Wide, parallelizable work: codebase-wide bug hunts · migrations touching hundreds of files · security audits · performance profiling across services.

NOT FOR
A faster way to write a single function. Interactive, ambiguous work stays in the normal conversation loop.

A workflow that modifies code is only as safe as your test gates — it can run the tests as part of the process, but only if the tests exist.

Best argument yet for writing tests first. · Research preview since May · CLI, Desktop, VS Code · Max / Team / Enterprise.

GUARDRAILS & WIRING

Four knobs worth knowing.

KNOB	SETTING	WHY YOU CARE
`size`	**Default guideline: “medium” — < 15 agents** `/config → Dynamic workflow size`	Wider isn't better. 15 is the line where coordination cost starts eating the win.
`nesting`	**Subagents nest to depth 3 by default** `CLAUDE_CODE_MAX_SUBAGENT_SPAWN_DEPTH=1`	Set it to 1 to disable nesting entirely while you're learning the shape of a workflow.
`routing`	**Bill the fan-out at mid-tier** `CLAUDE_CODE_SUBAGENT_MODEL=claude-sonnet-5`	Opus 5 does the decomposition; every delegation bills Sonnet. One env var, permanent saving.
`effort`	**Composes per node** `low → mechanical nodes · xhigh → judge / verify`	The dial isn't one global setting — it's a property of each node in the graph.

That last pair is the map's two zones, wired together — fan-out priced per attempt, judgment priced per mistake.

TWO STORIES THAT SELL IT

Nobody had to ask twice.

THE BUN REWRITE · VERIFIED

Zig → Rust, in 11 days.

535,496	lines of Zig converted, across 1,448 files
May 3–14	2026 — 11 days, 6,502 commits
64	Claude instances in parallel, at peak

One prompt plus steering, with adversarial Claude review catching critical bugs pre-merge. The enabling mechanism: **massive test suites on both sides, acting as the verifier**.

A task previously scoped in quarters.

COMMUNITY DEMO · UNPROMPTED

27 agents. One per part.

Someone sent Opus 5 to build an F1 showroom in Blender.

It drafted the car — then fired off workflows with **27 agents, one per part**, each building a high-detail piece.

Nobody asked it to. **The decomposition instinct is in the model** — Dynamic Workflows just gives it somewhere to put it.

Bun figures fact-checked to the official Bun engineering blog · bun.com/blog/bun-in-rust

LIVE DEMO · ONE TOOL, THREE ACTS

The dial, then a redline round.

0
The dial
Terminal output of **effort_sweep**: the same prompt on Opus 5 at **low** vs **xhigh**, token counts side by side.

PRE-RUN · 60 SECONDS

1
Judgment + collaboration
A volunteer marks up a draft: tighten this, fact-check that, a question, an arrow. One round back — surgical edits, a **reply** to the question, a correction marked *proposed*.

OPUS 5 · XHIGH

2
Smart routing
While Act 1 runs: the three routing layers, cheapest decision-maker first.

NEXT SLIDE

3
The workflow
Six sources, **one agent per source**, then a cross-check that drops anything only one source carries. Open the generated script: **haiku** on the fan-out, **opus** on the judge.

THE PAYOFF

It's a collaborator, not a compiler — the question got a reply, not an edit.

Prime directive: the author's words are sacred. Unannotated paragraphs are untouched.

ACT 2 · THE THREE ROUTING LAYERS

Cheapest decision-maker first.

1. **Static routing** — config

   `CLAUDE_CODE_SUBAGENT_MODEL=claude-sonnet-5`

   Every delegation bills mid-tier by default. One env var, permanent saving, zero decisions at runtime.

2. **Policy routing** — agent definitions

   `.claude/agents/explorer.md` → `model: haiku`, tools Read/Grep/Glob only

   “Return conclusions, not file dumps.” The policy is written once by a human or by Opus. Haiku just executes it.

3. **Judgment routing** — the model itself

   “Before changing anything, survey this page's styling system — layout, spacing scale, color tokens — conclusions only.”

   Opus 5 fans the survey out to the explorer instead of reading everything itself. Your conversation stays clean.

> **Routing gets smarter as you go down the list, and more expensive — so push every decision as far up as it can live.**

ACT 3 OUTPUT · RUN 2026-07-31

Six claims went in. Two came back.

CLAIM	WHAT CHECKING FOUND	VERDICT
$5 / $25 price	Resellers publish the same number because they *resell at it*. Propagation, not corroboration.	CONFIRMED
July 24 ship date	An administrative fact — an event Anthropic schedules. Independently reported.	CONFIRMED
Frontier-Bench 43.3 “new SOTA”	Every instance traces to one internal Anthropic run; the maintainers haven't published the dataset or the code. Epoch AI — genuinely independent — has Opus 5 and Fable 5 effectively tied.	SELF-REPORTED
CursorBench near-parity	Cursor alone builds, runs, grades and publishes it. BenchLM excludes it from scoring for exactly that reason.	SINGLE-CHANNEL
Fast mode ~2.5× speed	The price is verified live. No independent throughput measurement exists — a hole where a test should be.	UNMEASURED
ARC-AGI-3 30.2%	The only claim with an independent administrator — and the only one that drew a live dispute. OpenAI's 2026-07-30 post: 38.3% with two settings.	DISPUTED

The survivors aren't the *true* claims — they're the administrative ones.

Contradiction is a marker of instrumentation. Read silence as a warning, not a clearance.

TAKE THREE THINGS

01 Make Opus 5 the daily driver.

Tune cost with **effort**, not by downgrading models.

02 Update your prompts by *removing* things.

Verification scaffolding now costs money and adds nothing.

03 Route deliberately.

Opus for judgment, Sonnet for implementation, Haiku for fan-out — and reach for a Dynamic Workflow when the work is wide, parallel, and test-gated.

WHERE TO READ THE REST

Links.

Opus 5 announcement
anthropic.com/news → “Introducing Claude Opus 5”

Opus 5 prompting guide
docs.claude.com → “Claude Opus 5 prompting best practices”

Dynamic Workflows
docs.claude.com → “Claude Code: Dynamic Workflows (research preview)”

The Bun rewrite receipt
bun.com/blog/bun-in-rust

Sonnet 5 announcement
anthropic.com/news → “Introducing Claude Sonnet 5”

Effort docs
docs.claude.com → “Effort control · output_config.effort”

Intelligence vs Cost chart
artificialanalysis.ai → “Intelligence Index vs Cost per Task”

This deck & the map
talk/slides.html · talk/opus-5-cost-capability-map.html

Thank you. Questions — I have about fifteen minutes.